

## CS 425 MP2 Report

### GA10 - Sunveg Nalwar (snalwar2) and Sudarshan Shinde (sshinde5)

#### Design:

We implement a full-membership protocol with two variants that share a common SuspicionManager. Gossip sends periodic UPDATE\_BATCH heartbeats to  $k$  random peers and, with suspicion enabled, promotes ALIVE->SUSPECT after  $T_{fail}$  of silence and SUSPECT->DEAD after  $T_{cleanup}$ , disseminating updates via gossip and piggyback. PingAck actively probes one peer per period (PING/ACK) and piggybacks membership updates; with suspicion on, missed ACKs create SUSPECT and later DEAD; with suspicion off, we use per-target no-ACK timers to mark DEAD. Joins disseminate via an introducer JOIN\_ACK snapshot plus a one-shot push; leaves are explicit LEFT with higher incarnation and immediate fanout. To prevent cold-start storms when turning suspicion on, we initialize last-heard for unseen peers and apply a brief grace window. DEAD entries are Garbage collected after a timeout. Bandwidth is minimized by compact protobuf batches and piggyback in PingAck. We tune  $T_{fail}/T_{cleanup}$  and tick periods to satisfy the 3s/6s completeness bound while keeping background bytes/sec low.

#### **Protocol period calculation:**

Settings we use in experiments: period=300 ms (PING and gossip tick),  $T_{fail}=1$  s,  $T_{cleanup}=1$  s, FanoutK=3, GC TTL=5 s, suspicion grace=2 s only at switch-time.

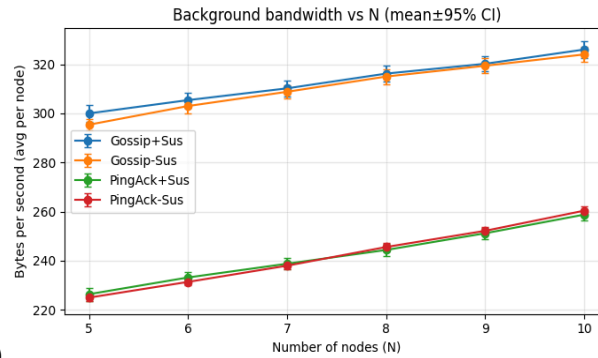
PingAck+Suspicion: A failure is detected when a probed node misses an ACK and remains silent for  $T_{cleanup}$ . Worst-case time to first SUSPECT is  $\leq$  one probe cycle to pick the failed node ( $\leq$  period) plus up to  $T_{fail}$  of not hearing (in practice, we probe before  $T_{fail}$ , but bound with period+ $T_{fail}$ ), and promotion to DEAD after  $T_{cleanup}$ . So first confirmed failure (DEAD)  $\leq$  period +  $T_{fail}$  +  $T_{cleanup} \approx 0.3 + 1 + 1 = 2.3$  s. This meets the 3 s bound.

Gossip+Suspicion: Detection is silence-based. First SUSPECT  $\leq T_{fail}$  (since we check every period), then DEAD after  $T_{cleanup}$ . So first confirmed failure (DEAD)  $\leq T_{fail} + T_{cleanup} \approx 2$  s. Dissemination bound: With FanoutK=1 (or piggyback-only in ping), worst-case spread is  $O(N)$  periods in adversarial order. Using  $N=10$  and period=0.3 s gives  $\leq 3$  s to touch all nodes once. In practice we have both piggyback and (optionally) one-shot fanout on joins/DEAD, so observed convergence is typically  $\leq 2-3$  s. Overall worst-case from failure to all tables updated:  $\leq$  (first DEAD) + (spread)  $\approx 2.3 + 3.0 \approx 5.3$  s (ping+suspect), and  $\approx 2.0 + 3.0 \approx 5.0$  s (gossip+suspect), satisfying the 6 s completeness requirement. If you run FanoutK=2-3, convergence is faster and less bursty.

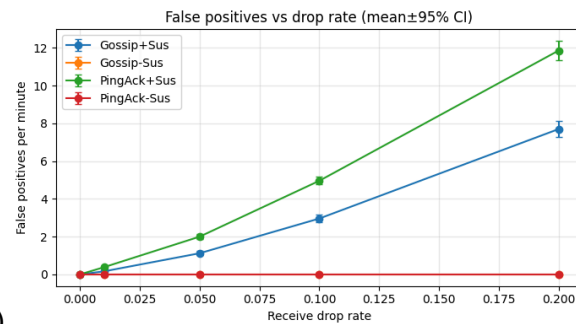
#### **False positives introduction and calculation:**

We drop any inbound UDP (including UPDATE\_BATCH and ACK) before the handler, simulating loss/delay as required. This is done via a command on CLI - drop [0,1]


A false positive is any SUSPECT that is later refuted locally without becoming DEAD. Our code logs "SUSPECT <node> ..." and "REFUTE <node> (heard again)"; we count a FP when a SUSPECT is followed by REFUTE for the same node and incarnation without an intervening DEAD. Rate reported as FPs/minute over a run (recommended 5-10 minutes per point)



Bandwidth (1a)



False positives (1b)



Detection time (1c)

2d) PingAck benefits more. Suspicion converts a missed ACK into immediate SUSPECT (then DEAD after Tcleanup), avoiding waiting for long silence on a peer that might otherwise not be probed again quickly. Gossip already checks silence periodically for every peer; suspicion mainly shortens the confirmation window, so the marginal gain is smaller than in PingAck.

2e) Gossip-Sus generally detects and converges faster than Ping-Sus because periodic heartbeats/gossip give consistent anti-entropy to all peers, while Ping-Sus relies on piggyback and chance probe paths, making it more sensitive to loss and target selection.

